

Prompter Use Cases and Tools

What to build next on Prompter using NVIDIA free-endpoint models from build.nvidia.com

DOWNLOAD

TRANSCRIBE

TRANSFORM

38 FREE MODELS VERIFIED 2026-10-04

Prompter today - verified from the live page

Prompter at prompter.neuralforge.tech is an AI media workspace with no sign-up. Two tools are live: **(1) Download video** - paste a media link, choose Best / 1080p / 720p / 480p / 360p (YouTube is explicitly not supported because it blocks datacenter IPs); and **(2) Create transcript** - speech to clean TXT + SRT, optional speaker diarization (Speaker 1, 2... up to 20), email-me-the-result, shareable job links, results kept a few days. Jobs run via GitHub Actions + Cloudflare Worker with queued / running / done states.

Two items are already on the site roadmap as *Coming soon*: **Translate across languages** (preserve meaning, structure and speaker context) and **Publish-ready subtitles** (styled captions burned onto video). This guide treats those as **already-planned - accelerate them**, not as new discoveries.

Strategy in one line: Keep the download -> transcribe core exactly as it is, and add an **AI transform layer on top of every transcript**. NVIDIA free endpoints do not replace Prompter's transcription engine - they turn each finished transcript into summaries, translations, searchable knowledge, repurposed content and new media.

What this guide gives you

- **15 tools (13 recommended + 2 experimental)** in 3 tiers - each as a card with the problem, the Prompter user flow, the exact NVIDIA model id, why that model, build effort and fit priority.
- A **persona x tool matrix** (creator, podcaster, teacher/student, journalist, marketer, accessibility user, fitness creator).
- A **model reference table** and a **Phase 1 / 2 / 3 build order** with dependencies.
- A **risks and verify-first checklist** - including an unverified Pakistan phone-verification flag you must test before committing.

All NVIDIA figures are **as shown on build.nvidia.com on 2026-10-04 - re-verify before building**. Model ids are renamed and deprecated without notice. No pricing is quoted because the free tier has no per-token billing to quote - see the facts box on the next page.

How to read this - NVIDIA free-tier facts

Read this box before any card. Every tool in this guide assumes these conditions, and every condition can change without notice.

Free-endpoint reality (verified 2026-10-04 by live browsing build.nvidia.com/models):

- **97 models total** in the catalog; filtering "**Free Endpoint**" = **38 models**. All 38 are enumerated in the Appendix.
- **Chat models** call <https://integrate.api.nvidia.com/v1> with \$NVIDIA_API_KEY (OpenAI-compatible). Get a key from a free NVIDIA Developer Program account; no credit card.
- NVIDIA Help FAQ: "**up to 40 requests per minute (RPM) for most models, with no per-token billing**"; all models are free to prototype. Treat 40 RPM as a ceiling under shared traffic, not a guarantee - queue jobs and cache results.
- **NVIDIA may record input/output** to provide the trial and improve its products. Add a user disclosure on Prompter and avoid promising confidentiality for sensitive media.
- **Specialized models are NOT OpenAI chat**: Kumo uses ai.api.nvidia.com/v1/structured-data/ . . . ; 3D Body Pose uses gRPC grpc.nvcf.nvidia.com:443. Several free models are **Experience-demo only** (try in browser, no API).
- **Model ids get renamed/deprecated without notice**. Pin ids in config, monitor the catalog, and never hard-code an id you have not re-checked. Example in this catalog: `riva-translate v1_1` is badged "Deprecation in 1d" - do not build on it.
- **21 of 38 are "Downloadable + Free"** - you can self-host the NIM container if the hosted endpoint is unavailable to you.

All figures above: as shown on build.nvidia.com on 2026-10-04 - re-verify before building.

The gaps - what the free tier does NOT give you

- **No standalone free ASR model**. Parakeet is in the wider catalog but is **not** in the 38 free endpoints. NVIDIA cannot replace Prompter's current transcription engine; it layers on top of finished transcripts. (See Risks.)
- **No free reranker**. Only one embedding model is free (`nvidia/nemotron-3-embed-1b`). Design search without a rerank stage.
- **No standalone Free Endpoint OCR and no image-generation/editing models** in the 38. `nemotron-parse / ocr-v2` exist in the full catalog but are **not free** - do not promise OCR or image-gen tools on the free tier.
- **Speech caveat**: `magpie-tts-zeroshot` and `nemotron-voicechat` need access approval despite the Free badge.

How each tool card is scored

Build effort: Low = prompt + one API call on an existing transcript, no new infra. Med = new pipeline stage, chunking, storage or a specialized endpoint. High = new transport (gRPC), access approval, or heavy media processing. **Fit priority**: how directly the tool extends Prompter's download -> transcribe -> transform promise for its existing audience.

Tier 1 - Build first: direct pipeline extensions

These six tools need no new media engine. Each one takes the transcript Prompter already produces and returns something more valuable than raw text. Ship in this order.

1. AI Summary + Chapters + Key Moments FIT: HIGHEST EFFORT: LOW

Problem: A 60-minute transcript is searchable but not skimmable. Users re-read or abandon it.

Flow: paste link -> transcribe (existing) -> **new: "Summarize"** -> TL;DR, chapter list with timestamps, 5-10 key moments with quotes -> export TXT / shareable link.

Model: nvidia/nemotron-3-super-120b-a12b (primary, 1M context handles long transcripts) or glm-5-3-flash (faster, multimodal fallback). Chat endpoint. **Why:** long-context reasoning models segment and summarise in one pass; no chunk-merge complexity for most videos.

2. Transcript Translator ALREADY-PLANNED - ACCELERATE FIT: HIGHEST EFFORT: LOW-MED

Problem: This is the site's own "Translate across languages" roadmap item. Users with non-English media are blocked today.

Flow: transcribe -> pick target language (37 supported) -> translated TXT + SRT with speaker labels preserved -> export.

Models: nvidia/riva-translate-4b-instruct-v2 (37 languages, few-shot prompts for names/terms) for the translation pass, then a light LLM polish pass (nemotron-3.5-lightning-30b-a3b) only where tone needs fixing. **Why:** Riva is purpose-built for translation; do not waste a large LLM on the bulk pass. **Do not use** riva-translate v1_1 (12 languages, "Deprecation in 1d" on 2026-10-04).

3. Semantic Search & Chat-with-Video FIT: HIGH EFFORT: MED

Problem: Keyword search misses paraphrase. Users ask "where did they explain pricing?" and get nothing.

Flow: transcribe -> transcript auto-chunked and embedded in background -> user searches by meaning or asks a question -> answer with timestamped citations jumping to the moment.

Models: nvidia/nemotron-3-embed-1b (1B embedding, query/passage modes, built for RAG) for indexing + any Tier-1 LLM for the answer. **Why:** this is the only free embedding model, and it unlocks Tools 3, 7 and 13. Build the embedding store once, reuse everywhere. No free reranker - rank by cosine similarity alone.

Tier 1 (continued)

4. Content Repurposer FIT: HIGH EFFORT: LOW

Problem: Creators transcribe a video, then manually rewrite it into every other format. That rewriting is the actual job.

Flow: transcribe -> **"Repurpose"** -> pick outputs: blog post, YouTube description with timestamps, X thread, LinkedIn post, newsletter section -> edit and export.

Model: kimi - k3 (long-horizon, strong long-form writing) or nemotron-3-ultra-550b-a55b for the highest-quality pass. Chat endpoint. **Why:** one transcript, templated prompts per format. Highest perceived value per API call of any card here.

5. Smart Subtitles ALREADY-PLANNED - ACCELERATE FIT: HIGHEST EFFORT: MED

Problem: The site's second roadmap item - "Publish-ready subtitles, styled captions burned onto video." Raw SRT is not publish-ready: line breaks are awkward and timing is word-dump, not reading rhythm.

Flow: transcribe (existing SRT) -> LLM re-segments into readable caption lines respecting phrase boundaries -> style picker -> burned-in video export (existing pipeline) + clean SRT download.

Model: nemotron-3.5-lightning-30b-a3b or glm-5-3-flash for fast re-segmentation. **Why:** this is a language task, not a vision task - a fast LLM fixing breaks and punctuation is the whole upgrade. Pairs with Tool 2 for translated burned-in subtitles.

6. Meeting / Podcast Notes FIT: HIGH EFFORT: LOW

Problem: Speaker-diarized transcripts already exist on Prompter, but nobody wants to read 90 minutes to find the three decisions.

Flow: upload/paste meeting or podcast -> transcribe with diarization -> **"Notes"** -> action items by speaker, decisions, open questions, pull-quotes with timestamps -> export / email.

Model: nvidia/nemotron-3-super-120b-a12b (structured extraction, tool calling for JSON action items). **Why:** diarization is already a Prompter strength - this monetises it. Structured JSON output makes the result exportable to Notion/Docs.

Tier 1 sequencing note: Build Tool 1 first (proves the LLM layer), Tool 3's embedding store second (unlocks search, chat and clips), then 2, 4, 5, 6 in any order - they are independent prompts on shared infra.

Tier 2 - New tools, strong media fit

These open new audiences rather than deepening the current one. Each needs a specialized model or a new media stage, so effort is higher - but the media fit is direct.

7. Translated Dubbing FIT: HIGH EFFORT: MED-HIGH

Problem: Subtitles localise text; creators want their own voice speaking another language.

Flow: transcribe -> translate (Tool 2) -> voice-clone TTS from a short sample of the original speaker -> dubbed audio track muxed onto the downloaded video.

Models: nvidia/riva-translate-4b-instruct-v2 + magpie-tts-zeroshot (voice-clone TTS from a short sample, 23 languages). **Caveat:** Magpie shows "Apply for Access" despite the Free badge (2026-10-04) - apply early, build the pipeline against the API shape meanwhile.

8. Audio Cleanup for Podcasters FIT: HIGH EFFORT: MED

Problem: Bad room noise ruins otherwise good recordings; podcasters currently leave Prompter to fix audio elsewhere.

Flow: upload audio/video -> "Clean audio" -> noise-removed, studio-voice enhanced track -> download cleaned file + transcript of the cleaned audio.

Models: Background Noise Removal (bnr) + Studio Voice (mic enhancement), both in the 38 free endpoints. Specialized (non-chat) endpoints. **Why:** a pre-transcription cleanup pass also improves transcript accuracy - two wins from one stage.

9. Video Describer / Video Q&A FIT: HIGH EFFORT: MED

Problem: Transcripts capture what was said, not what was shown. Tutorials, demos and reviews live in the visuals.

Flow: paste link -> frames sampled + transcript aligned -> scene-by-scene visual description, or ask "what does he click at 12:40?" -> timestamped answers.

Models: nvidia/nemotron-3-nano-omni-30b-a3b-reasoning (document/image/video/audio, OCR-labelled) or cosmos3-nano-reasoner (physical-world video reasoning); meta/llama-3.2-90b-vision-instruct for frame-level questions. Chat/vision endpoint.

Tier 2 (continued) and Tier 3

10. Image Describer & Alt-Text Generator **FIT: MED-HIGH** **EFFORT: LOW**

Problem: Accessibility users and creators need descriptions of thumbnails, slides and screenshots - today there is no image tool on Prompter at all.

Flow: upload image (or auto-extract video thumbnail/slides) -> short alt-text, long description, and extracted on-image text summary -> copy/export.

Models: meta/llama-3.2-11b-vision-instruct (fast) or paligemma. **Honest limit:** there is no standalone free OCR model - this describes and summarises images; it does not promise verbatim OCR extraction.

11. AI-Video Detector **FIT: MED-HIGH** **EFFORT: LOW-MED**

Problem: Journalists and teachers need a first check: is this clip AI-generated before they cite or share it?

Flow: paste link / upload clip -> "**Check authenticity**" -> synthetic-probability report with a plain-language explainer and a "not proof - verify source" disclaimer.

Model: synthetic-video-detector (AI-generated video detection, in the 38). Specialized endpoint. **Why:** a trust tool differentiates Prompter from every downloader clone. Never present the score as proof.

12. Content Safety Check **FIT: MED** **EFFORT: LOW**

Problem: Brands and teachers publishing repurposed content need a toxicity/unsafe-content pass before it goes out.

Flow: transcript (or repurposed draft) -> "**Safety check**" -> flagged segments with category and severity -> edit and re-check.

Models: nemotron-3.5-content-safety (multilingual, multimodal) or llama-3.1-nemotron-safety-guard-8b-v3 / llama-guard-4-12b. Chat/safety endpoint. Also useful as Prompter's own moderation layer for public shareable links.

13. Shorts / Clip Finder **FIT: HIGH** **EFFORT: MED**

Problem: The most valuable 30 seconds of a long video are buried. Creators scrub manually for Shorts/Reels material.

Flow: transcribe + embed (Tool 3 store) -> segments scored for hook, self-containedness and quotability -> ranked clip list with timestamps -> jump, preview, export ranges.

Models: nvidia/nemotron-3-embed-1b (segment embeddings) + LLM scoring (glm-5-3-flash). **Why:** reuses Tier-1 infrastructure; the output feeds directly into the download tool's quality selector.

Tier 3 - Experimental / niche (build only with a named audience)

14. 3D Pose Tracker for Fitness & Dance Creators EFFORT: HIGH FIT: NICHE

Model: 3D Body Pose (video -> 3D skeleton). **Costs:** gRPC transport (grpc.nvcf.nvidia.com:443), 50MB H.264 mp4 input limit. High effort, small audience - build only if fitness/dance creators ask by name.

15. Live Voice Translator / Assistant EFFORT: HIGH FIT: NICHE

Model: nemotron-voicechat (ASR+LLM+TTS full-duplex). **Caveat:** "Apply for Early Access" on 2026-10-04. A live product is a different company from an async job queue - park until access is granted and demand is proven.

Explicitly NOT recommended for Prompter: Kumo Relational (multi-table structured-data prediction - enterprise data-science audience, non-media, specialized endpoint), **ising-calibration 1.5 / 1.0** (quantum-computing calibration plots), and **streampetr / sparsedrive / bevformer** (autonomous-driving perception). All are in the 38 free endpoints; none fits Prompter's media audience. Listing them here so the catalog feels fully triaged, not cherry-picked.

Use-case matrix - who each tool serves

Read rows to prioritise by audience; read columns to see which tools serve everyone. Tier-1 tools (1-6) serve nearly every persona - that is why they ship first.

Persona	Start here	Then	The job it finishes
Creator (YouTube/TikTok)	1 Summary, 4 Repurposer, 13 Clip Finder	5 Smart Subtitles, 7 Dubbing	One long video becomes a week's content
Podcaster	8 Audio Cleanup, 6 Podcast Notes	1 Summary, 4 Repurposer	Clean audio, show notes, quotes - no manual pass
Teacher / Student	3 Chat-with-Video, 1 Summary	2 Translator, 12 Safety Check	Lectures become searchable, citable study material
Journalist	11 AI-Video Detector, 3 Semantic Search	6 Notes, 2 Translator	Verify a clip, then find the exact quote in minutes
Marketer	4 Repurposer, 12 Safety Check	13 Clip Finder, 7 Dubbing	Webinars become brand-safe, localised campaigns
Accessibility user	5 Smart Subtitles, 10 Alt-Text	9 Video Descriptor, 2 Translator	Media that can be read, heard and understood
Fitness creator	13 Clip Finder	14 Pose Tracker (Tier 3)	Form-check clips without manual scrubbing

Founder read: If you can only market to one persona at launch, market Tier 1 to **creators and podcasters** - Tools 1, 4, 6 and 13 form a single story ("paste a link, get a week of content") and every one runs on infrastructure you build once.

What not to promise

- Do not promise **replacement transcription** ("faster NVIDIA transcription") - no free ASR exists to back it.
- Do not promise **OCR or AI image generation** - no free endpoint in the 38 supports either.
- Do not promise **instant** results on the free tier - 40 RPM shared traffic means queues; Prompter's existing queued/running/done UX is the honest pattern.

Model reference - the free models Prompter would actually call

Subset of the 38 free endpoints with a Prompter use. Full grouped list of all 38 in the Appendix. All rows: as shown on build.nvidia.com on 2026-10-04 - re-verify before building.

Model id	Family	Prompter uses it for	Endpoint
nvidia/nemotron-3-super-120b-a12b	LLM	Summary, chapters, meeting notes (Tools 1, 6)	Chat
nvidia/nemotron-3-ultra-550b-a55b	LLM (1M ctx)	Highest-quality repurposing (Tool 4)	Chat
nvidia/nemotron-3.5-lightning-30b-a3b	LLM (fast)	Caption re-segmentation, polish (Tools 2, 5)	Chat
glm-5-3 / glm-5-3-flash	LLM (multimodal flash)	Fast summary fallback, clip scoring (Tools 1, 13)	Chat
kimi-k3	LLM (coding/agentic)	Long-form repurposing (Tool 4)	Chat
nvidia/riva-translate-4b-instruct-v2	Translation (37 lang)	Translator + dubbing script (Tools 2, 7)	Specialized
nvidia/nemotron-3-embed-1b	Embedding (only free one)	Semantic search, chat, clips (Tools 3, 13)	Specialized
nvidia/nemotron-3-nano-omni-30b-a3b-reasoning	Vision (doc/img/video/audio)	Video describer / Q&A (Tool 9)	Chat (vision)
cosmos3-nano-reasoner	Vision (video reasoning)	Physical-world video Q&A (Tool 9)	Specialized
meta/llama-3.2-90b-vision-instruct	Vision	Frame-level questions (Tool 9)	Chat (vision)
meta/llama-3.2-11b-vision-instruct / paligemma	Vision (fast)	Image describer, alt-text (Tool 10)	Chat (vision)
magpie-tts-zeroshot	Speech TTS (voice clone)	Translated dubbing (Tool 7)	Specialized*
nemotron-voicechat	Speech (full-duplex)	Live voice assistant (Tool 15)	Specialized*
Background Noise Removal (bnr) / Studio Voice	Speech (audio enhance)	Audio cleanup (Tool 8)	Specialized
synthetic-video-detector	Video (detection)	AI-video detector (Tool 11)	Specialized
nemotron-3.5-content-safety / llama-3.1-nemotron-safety-guard-8b-v3 / llama-guard-4-12b	Safety	Safety check (Tool 12)	Chat/safety

Model reference (continued)

Model id	Family	Prompter uses it for	Endpoint
Active Speaker Detection	Video	Possible diarization assist (future)	Specialized
3D Body Pose	Video (pose)	Pose tracker (Tool 14, Tier 3)	gRPC

* Magpie and VoiceChat show "Apply for Access / Early Access" despite the Free badge (2026-10-04). Models named without a Prompter use (Kumo, ising-calibration, autonomous-driving family, cosmos generation models, remaining chat models) are listed in the Appendix with the reason they are not in this table.

Build order roadmap

Phase	Ship	Depends on
Phase 1 - Transform layer	Tool 1 Summary/Chapters; Tool 3 embedding store + Semantic Search; Tool 4 Repurposer; Tool 6 Notes	One chat-endpoint integration (<code>integrate.api.nvidia.com/v1</code>), one embedding store. No new media processing.
Phase 2 - Roadmap + reach	Tool 2 Translator (accelerates Coming-soon item); Tool 5 Smart Subtitles (accelerates second item); Tool 13 Clip Finder; Tool 8 Audio Cleanup; Tool 12 Safety Check	Phase-1 embedding store (for 13); Riva integration (for 2, then 7); specialized audio endpoints (for 8).
Phase 3 - New surfaces	Tool 9 Video Describer/Q&A; Tool 10 Alt-Text; Tool 11 AI-Video Detector; Tool 7 Dubbing (once Magpie access granted)	Frame-sampling stage (for 9); Magpie access approval (for 7).
Parked	Tools 14-15 (Tier 3)	Named audience demand; VoiceChat early-access grant.

Prompter Use Cases and Tools

Risks & verify-first checklist

(a) TEST FIRST - account access from your location is UNVERIFIED. A third-party tracker (*awesome-free-ai-coding*, read 2026-09-25) lists **Pakistan among countries excluded from NVIDIA phone verification** for the free Developer key. NVIDIA's own pages opened for this guide showed **no country list**, so this is **unverified - do not treat it as settled either way**. Before committing to any tool above: create the NVIDIA Developer account and API key from your actual location and make one test call. If blocked, options are an eligible-region account operated lawfully, partner endpoints, or self-hosting the **Downloadable + Free NIMs** (21 of 38).

- **(b) 40 RPM throttling.** The FAQ figure is a shared-traffic ceiling. Prompter already queues (GitHub Actions + Cloudflare Worker) - keep every NVIDIA call inside that queue, cache by transcript hash, and never call synchronously from page load.
- **(c) NVIDIA may record inputs/outputs.** Add a plain disclosure ("media you submit may be processed by NVIDIA and recorded under its trial terms") and do not position Prompter for confidential media until you have a paid/self-hosted path.
- **(d) Model churn.** `riva-translate v1_1` deprecating in this very catalog is the warning. Pin ids in one config file, add a weekly catalog check, and design prompts to survive a model swap.
- **(e) The gaps are real.** No free ASR (Parakeet is not free), no free OCR, no free image generation. Do not let marketing copy promise any of them.
- **(f) Access-gated "free" speech models.** Magpie TTS and VoiceChat need approval despite the Free badge. Apply on day one of Phase 1 so the grant lands before Phase 3 needs it.
- **(g) Trial terms.** Free endpoints are for prototyping under NVIDIA's Trial Terms (18+). Read the current terms before a public launch and price the paid/self-hosted exit.

Appendix - all 38 free-endpoint models (2026-10-04)

Grouped as enumerated from build.nvidia.com/models with the "Free Endpoint" filter on 2026-10-04. Re-verify before building - ids are renamed/deprecated without notice.

LLM / reasoning / coding (12)

nvidia/nemotron-3-ultra-550b-a55b (1M ctx) · nvidia/nemotron-3-super-120b-a12b · nvidia/nemotron-3.5-lightning-30b-a3b · deepseek-v4.1-flash (multimodal) · glm-5-3 · glm-5-3-flash (multimodal) · kimi-k3 (Downloadable+Free) · muse-glimmer-30b (Meta, text+image) · diffusion-gemma-26b-a4b-it (Google) · laguna-xs-2.1 (Poolside) · openai/gpt-oss-20b · gemma-4-31b-it (Google)

Vision (5)

meta/llama-3.2-90b-vision-instruct · meta/llama-3.2-11b-vision-instruct · paligemma · nvidia/nemotron-3-nano-omni-30b-a3b-reasoning · cosmos3-nano-reasoner

Embeddings (1) · Translation (2) · Safety (3)

Embedding: nvidia/nemotron-3-embed-1b (no free reranker). Translation: nvidia/riva-translate-4b-instruct-v2 (use this) · riva-translate_v1_1 (deprecating - do not build on it). Safety: nemotron-3.5-content-safety · llama-3.1-nemotron-safety-guard-8b-v3 · llama-guard-4-12b.

Speech (4) · Video (4)

Speech: magpie-tts-zeroshot* · nemotron-voicechat* · Background Noise Removal (bnr) · Studio Voice. Video: cosmos-transfer2.5-2b · cosmos3-nano (generation) · synthetic-video-detector · Active Speaker Detection. (*Apply for Access / Early Access despite Free badge.)

Other (7) - triaged, not recommended for Prompter

3D Body Pose (Tier 3 only) · Kumo Relational (structured-data prediction) · ising-calibration 1.5 · ising-calibration 1.0 (quantum) · streampetr · sparsedrive · bevformer (autonomous driving).

Sources

- Prompter live page: <https://prompter.neuralforge.tech/> (read 2026-10-04)
- NVIDIA model catalog: <https://build.nvidia.com/models> (Free Endpoint filter = 38 models, read 2026-10-04)
- NVIDIA API docs: <https://docs.api.nvidia.com> (chat base <https://integrate.api.nvidia.com/v1>; Kumo and Body Pose specialized endpoints as noted)
- NVIDIA Help FAQ (40 RPM, no per-token billing, input/output recording) - via build.nvidia.com, read 2026-10-04
- Third-party free-tier tracker flagging Pakistan phone-verification exclusion: *awesome-free-ai-coding*, read 2026-09-25 - **UNVERIFIED against NVIDIA's own pages; test before relying on it**

Final word for the builder: Prompter's moat is the queue, the diarized transcript and the shareable link - infrastructure most NVIDIA demos lack. The free tier gives you the intelligence to put on top of it. Ship Phase 1, measure which transform users re-run, and let that - not the catalog - choose Phase 2.